

✓ Project Report: Unconditional Generative Adversarial Network (GAN) on MNIST

Objective

The objective of this project is to implement an **Unconditional Generative Adversarial Network (GAN)** for the MNIST dataset. The goal is to train a generator capable of producing realistic handwritten digit images through adversarial training against a discriminator.

1. Environment and Data Setup

The model is implemented using **PyTorch** and optimized for GPU acceleration when available.

- **Dataset:** MNIST handwritten digits (60,000 images)
 - **Image Size:** 28×28 grayscale
 - **Normalization:** Images scaled to the range $[-1, 1]$
 - **Batch Size:** 128
-

2. Model Architecture

Generator

The Generator maps a latent noise vector sampled from a standard normal distribution to a flattened image representation.

- Input: 100-dimensional noise vector
- Fully connected layers with ReLU activation
- Output layer with Tanh activation

Discriminator

The Discriminator distinguishes real MNIST images from generated images.

- Input: Flattened image vector
 - Fully connected layers with LeakyReLU activation
 - Output probability using Sigmoid activation
-

3. Loss Function and Optimization

The adversarial training objective is defined using **Binary Cross-Entropy (BCE) loss**.

- **Generator Objective:** Fool the discriminator
- **Discriminator Objective:** Correctly classify real and fake images

Optimizer: Adam

Learning Rate: 0.0002

Beta Values: (0.5, 0.999)

4. Training Procedure

Training follows an alternating optimization scheme:

1. The discriminator is trained on real and generated images.
2. The generator is updated to maximize discriminator error.

This adversarial process continues until convergence.

5. Evaluation and Analysis

The quality of generated images improves progressively during training. After sufficient epochs, the generator produces visually recognizable handwritten digits. Training remains stable without significant mode collapse.

6. Conclusion

This project demonstrates a successful implementation of an unconditional GAN on the MNIST dataset. The trained generator learns the underlying data distribution and generates realistic digit samples.

Code:

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
from tqdm import tqdm
import matplotlib.pyplot as plt

# --- 1. Environment & Hyperparameters ---
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
latent_dim = 100
image_dim = 28 * 28
batch_size = 128
lr_G = 2e-4
lr_D = 1e-4
epochs = 30

# --- 2. Generator & Discriminator (Deliverable A) ---
class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(latent_dim, 256),
            nn.BatchNorm1d(256),
            nn.ReLU(),

            nn.Linear(256, 512),
            nn.BatchNorm1d(512),
            nn.ReLU(),

            nn.Linear(512, 1024),
            nn.BatchNorm1d(1024),
            nn.ReLU(),

            nn.Linear(1024, image_dim),
            nn.Tanh()
        )

    def forward(self, z):
        return self.net(z)

class Discriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(image_dim, 1024),
            nn.LeakyReLU(0.2),

            nn.Linear(1024, 512),
            nn.LeakyReLU(0.2),

            nn.Linear(512, 256),
            nn.LeakyReLU(0.2),

            nn.Linear(256, 1)  # logits (no sigmoid)
        )

    def forward(self, x):
        return self.net(x)

# --- 3. Main Loop ---
def main():
    tf = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.5,), (0.5,))
    ])

    loader = DataLoader(
        torchvision.datasets.MNIST(
            './data',
            train=True,
            download=True,
            transform=tf
        ),
```

```

        batch_size=batch_size,
        shuffle=True
    )

    G = Generator().to(device)
    D = Discriminator().to(device)

    criterion = nn.BCEWithLogitsLoss()
    opt_G = torch.optim.Adam(G.parameters(), lr=lr_G, betas=(0.5, 0.999))
    opt_D = torch.optim.Adam(D.parameters(), lr=lr_D, betas=(0.5, 0.999))

    # Training
    print(f"\nTraining Unconditional GAN on {device}...")
    G.train(); D.train()

    for epoch in range(epochs):
        pbar = tqdm(loader, desc=f"Epoch {epoch+1}/{epochs}")
        for imgs, _ in pbar:
            imgs = imgs.view(-1, image_dim).to(device)
            bs = imgs.size(0)

            # Label smoothing
            real_labels = torch.full((bs, 1), 0.9, device=device)
            fake_labels = torch.zeros(bs, 1, device=device)

            # -----
            # Train Discriminator
            # -----
            z = torch.randn(bs, latent_dim, device=device)
            fake_imgs = G(z)

            d_loss = criterion(D(imgs), real_labels) + \
                criterion(D(fake_imgs.detach()), fake_labels)

            opt_D.zero_grad()
            d_loss.backward()
            opt_D.step()

            # -----
            # Train Generator
            # -----
            g_loss = criterion(D(fake_imgs), real_labels)

            opt_G.zero_grad()
            g_loss.backward()
            opt_G.step()

            pbar.set_postfix(D_Loss=d_loss.item(), G_Loss=g_loss.item())

    # --- 4. Generation (Deliverable B) ---
    print("\nGenerating realistic digits...")
    G.eval()
    with torch.no_grad():
        z = torch.randn(16, latent_dim, device=device)
        samples = G(z).view(-1, 1, 28, 28)

    # --- 5. Report (Deliverable C & D) ---
    print("\n" + "="*50 + "\nDELIVERABLES REPORT\n" + "="*50)
    print(f"(a) GAN Trained: YES (Final D Loss: {d_loss.item():.4f})")
    print(f"(b) Generation: COMPLETED (Samples below)")
    print(f"(c) Metrics: Final Generator Loss {g_loss.item():.4f}")
    print("="*50)

    grid = torchvision.utils.make_grid(samples.cpu(), nrow=4, normalize=True)
    plt.imshow(grid.permute(1, 2, 0))
    plt.axis('off')
    plt.show()

if __name__ == "__main__":
    main()

```

Training Unconditional GAN on cuda...

```
Epoch 1/30: 100%|██████████| 469/469 [00:16<00:00, 28.94it/s, D_Loss=0.944, G_Loss=1.32]
Epoch 2/30: 100%|██████████| 469/469 [00:15<00:00, 31.08it/s, D_Loss=1.23, G_Loss=1.49]
Epoch 3/30: 100%|██████████| 469/469 [00:15<00:00, 30.29it/s, D_Loss=1.06, G_Loss=1.2]
Epoch 4/30: 100%|██████████| 469/469 [00:15<00:00, 30.27it/s, D_Loss=1.06, G_Loss=1.41]
Epoch 5/30: 100%|██████████| 469/469 [00:15<00:00, 30.19it/s, D_Loss=1.05, G_Loss=1.4]
Epoch 6/30: 100%|██████████| 469/469 [00:15<00:00, 30.72it/s, D_Loss=1.38, G_Loss=0.702]
Epoch 7/30: 100%|██████████| 469/469 [00:15<00:00, 29.94it/s, D_Loss=1.17, G_Loss=0.912]
Epoch 8/30: 100%|██████████| 469/469 [00:15<00:00, 30.95it/s, D_Loss=1.14, G_Loss=1.25]
Epoch 9/30: 100%|██████████| 469/469 [00:15<00:00, 30.89it/s, D_Loss=1.2, G_Loss=0.863]
Epoch 10/30: 100%|██████████| 469/469 [00:15<00:00, 30.28it/s, D_Loss=1.16, G_Loss=1.09]
Epoch 11/30: 100%|██████████| 469/469 [00:15<00:00, 29.77it/s, D_Loss=1.24, G_Loss=0.712]
Epoch 12/30: 100%|██████████| 469/469 [00:16<00:00, 28.17it/s, D_Loss=1.22, G_Loss=1.4]
Epoch 13/30: 100%|██████████| 469/469 [00:17<00:00, 27.25it/s, D_Loss=1.24, G_Loss=0.957]
Epoch 14/30: 100%|██████████| 469/469 [00:15<00:00, 29.68it/s, D_Loss=1.22, G_Loss=1.13]
Epoch 15/30: 100%|██████████| 469/469 [00:15<00:00, 30.81it/s, D_Loss=1.33, G_Loss=1.03]
Epoch 16/30: 100%|██████████| 469/469 [00:15<00:00, 30.47it/s, D_Loss=1.29, G_Loss=0.923]
Epoch 17/30: 100%|██████████| 469/469 [00:16<00:00, 29.20it/s, D_Loss=1.2, G_Loss=1.19]
Epoch 18/30: 100%|██████████| 469/469 [00:15<00:00, 30.47it/s, D_Loss=1.2, G_Loss=1.05]
Epoch 19/30: 100%|██████████| 469/469 [00:15<00:00, 30.47it/s, D_Loss=1.18, G_Loss=1.26]
Epoch 20/30: 100%|██████████| 469/469 [00:15<00:00, 29.99it/s, D_Loss=1.39, G_Loss=1.34]
Epoch 21/30: 100%|██████████| 469/469 [00:15<00:00, 30.10it/s, D_Loss=1.27, G_Loss=1.18]
Epoch 22/30: 100%|██████████| 469/469 [00:15<00:00, 30.37it/s, D_Loss=1.31, G_Loss=1.26]
Epoch 23/30: 100%|██████████| 469/469 [00:15<00:00, 30.56it/s, D_Loss=1.27, G_Loss=1.26]
Epoch 24/30: 100%|██████████| 469/469 [00:16<00:00, 28.90it/s, D_Loss=1.23, G_Loss=0.952]
Epoch 25/30: 100%|██████████| 469/469 [00:15<00:00, 30.47it/s, D_Loss=1.3, G_Loss=0.87]
Epoch 26/30: 100%|██████████| 469/469 [00:15<00:00, 30.30it/s, D_Loss=1.34, G_Loss=0.64]
Epoch 27/30: 100%|██████████| 469/469 [00:15<00:00, 29.32it/s, D_Loss=1.22, G_Loss=1.25]
Epoch 28/30: 100%|██████████| 469/469 [00:15<00:00, 29.99it/s, D_Loss=1.29, G_Loss=1.1]
Epoch 29/30: 100%|██████████| 469/469 [00:15<00:00, 30.04it/s, D_Loss=1.3, G_Loss=1.52]
Epoch 30/30: 100%|██████████| 469/469 [00:15<00:00, 30.25it/s, D_Loss=1.23, G_Loss=0.939]
```

Generating realistic digits...

DELIVERABLES REPORT

- (a) GAN Trained: YES (Final D Loss: 1.2262)
- (b) Generation: COMPLETED (Samples below)
- (c) Metrics: Final Generator Loss 0.9393

